

Yuqing (Jacky) Jiang

Open to Relocate | (217)-898-8615 | jiangyuqing0508@outlook.com | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Illinois Urbana-Champaign

Bachelor of Science, Computer Science & Statistics (Dean's List)

Aug 2023 – May 2027

Cumulative GPA: 3.91/4.00

SKILLS & PUBLICATION

Core ML / DL: PyTorch, CUDA, Scikit-learn, XGBoost, LightGBM, Bayesian Optimization

LLM & Agentic Systems: LangChain, LangGraph, CrewAI, RAG, PEFT, Prompt Engineering, APIs (Claude, Azure OpenAI, Gemini)

System & Tools: Python, SQL, C++, Java, Docker, Azure, Git, CI/CD, TypeScript, Node.js, FastAPI, React, Linux

Publication: Steer2Adapt: Dynamically Composing Steering Vectors for Efficient LLM Adaptation — *ICLR 2026 Workshop (Accepted)* [arXiv:2602.07276]

WORK EXPERIENCE

Pacific Financial Professional Association | *Quantitative Finance AI Intern* | Atlanta, GA

Jun 2026 – Aug 2026

- Built a **4-layer market-data pipeline** (acquisition, cleaning, signal construction, publishing) in **Python/pandas** with versioned-Parquet inter-layer contracts, powering credit-risk analytics on a 150-company universe
- Implemented an **EM-calibrated Merton/KMV structural model** (asset value and volatility, distance-to-default, PIT/TTC PD) with a Time-Consistent (TiC) rating conversion, automating ticker to complete rating table in one command
- Quantified model risk with a 2,000-replicate bootstrap: the PD conversion layer amplified parameter uncertainty $\times 4,073$ relative to the drift-free risk score while rank ordering remained stable (**Kendall $\tau = 0.956$**); validated $\rho = 0.79$ against agency ratings, so every published rating ships with a confidence interval
- Enforced a **no-look-ahead timing protocol** (event/available/decision timestamps) for point-in-time correctness; shipped the production release with CI-gated tests ($\geq 85\%$ coverage) and a public interactive demo

Earnest Agriculture | *Marketing Analytics & Strategy Intern* | Champaign, IL

Jan 2026 – May 2026

- Boosted **organic traffic by 40% and sales by 14% YoY** by rewriting the blog into a Generative Engine Optimization (GEO) structure (featured-snippet summaries, FAQ blocks, entity lists) on Ahrefs keywords with 10k+ monthly searches and difficulty under 50
- Built an **LLM classification pipeline** (OpenAI API, prompt engineering) that labeled a 1,500+ post competitor dataset, scraped with Python/BeautifulSoup, by topic, format, and angle, replacing manual review
- Built a **RAG drafting workflow (LangChain, Chroma)** grounding blog drafts in indexed brand documents and Google Trends signals

SeeU International Inc. | *Data Analyst Intern* | Remote

Sep 2025 – Nov 2025

- Parsed 500K+ nested-JSON event records into **relational schemas** with Python and SQL, powering course-level KPI dashboards and funnel analytics
- Implemented NLP preprocessing and automated text-summarization workflows for user course feedback, improving data-driven curriculum iteration
- Built class-weighted churn-risk models (logistic regression, XGBoost, random forest; **AUC 0.78**) on platform-engagement features for the free-to-paid trial funnel, showing low-activity users accounted for 61% of churn and prioritizing them for re-engagement
- Designed an **A/B test** and **drove a 15% lift in student engagement** by translating the funnel and churn analysis into recommendations on course offerings and promotion timing

Walt Disney Company | *Operations & Analytics Intern* | Orlando, FL

Jun 2025 – Aug 2025

- Forecast demand for popcorn, bakery, and merchandise lines from POS and attendance data (30,000+ visitors, \$150,000+ daily sales) via attendance-driven regression with day-of-week effects against exponential-smoothing baselines, **reducing stockouts by 40%**
- Built inventory tracking dashboards monitoring sell-through and shrinkage to guide data-driven reorders and promotional displays

LLM RESEARCH

ScribeAR: Live Speaker Attribution for Accessible Captioning | *Research Assistant*

Feb 2026 – Present

- Built the live speaker-diarization feature for ScribeAR, UIUC's open-source real-time captioning suite for deaf, hard-of-hearing, and other users with disabilities: a **CPU-first pyannote pipeline** (community-1, CUDA-ready) running inside the whisper-streaming worker pool on a rolling 30-second buffer, so live captions carry stable speaker labels
- Designed a SpeakerReconciler that maps per-run labels to session-wide labels via **overlap voting** and assigns each word the speaker range it overlaps most; attached labels to the transcript sequence as an optional speakers array and never delay caption finalization
- Shipped the WCAG-AA React speaker display, with a colorblind-aware palette auto-adjusted to **4.5:1 contrast** so labels stay legible as speakers enter and leave, and a CPU/GPU benchmark harness (pyannote vs. Sortformer) for the 5-second tick; 30 unit tests

Steer2Adapt: Compositional Steering for LLM Adaptation | *Co-author, ICLR 2026 Workshop*

Sep 2025 – Jan 2026

- Built a training-free LLM adaptation framework in **PyTorch** that composes semantic steering vectors at inference time via Bayesian Optimization (Matern-5/2 GP, Expected Improvement), adapting to new tasks without gradient updates
- Achieved **+8.2% average accuracy across 9 tasks and 3 open-source LLMs** (Mistral, Qwen-2.5, Llama-3.1) from 12 calibration examples per task by building reusable semantic concept subspaces (Big Five traits for reasoning; refusal, sycophancy, hallucination, and fairness for safety)
- Designed a **stability-aware objective** that rewards fixing wrong predictions while penalizing flips of already-correct ones: gains on all 27 task-model pairs with no negative outliers, versus prompting baselines that regressed over 30% in some safety scenarios