

Yuqing (Jacky) Jiang

Open to Relocate | (217)-898-8615 | jiangyuqing0508@outlook.com | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Illinois Urbana-Champaign

Bachelor of Science, Computer Science & Statistics (Dean's List)

Aug 2023 – May 2027

Cumulative GPA: 3.91/4.00

SKILLS & PUBLICATION

Core ML / DL: PyTorch, CUDA, Scikit-learn, XGBoost, LightGBM

LLM & Agent Systems: PEFT, RAG, LangChain, CrewAI, Prompt Engineering, Steering Vectors

System & Tools: Python, C++, Java, Git, Docker, SQL, Linux, CUDA

Publication: Steer2Adapt: Dynamically Composing Steering Vectors for Efficient LLM Adaptation — *ICLR 2026 Workshop (Accepted)* / *NeurIPS 2026 (Under Review)* [arXiv:2602.07276]

WORK EXPERIENCE

Pacific Financial Professional Association | *Quantitative Finance AI Intern* | Atlanta, GA (Remote)

Jun 2026 – Present

- Built a **4-layer end-to-end market data pipeline** (acquisition, cleaning, signal construction, publishing) in **Python/pandas**, enforcing explicit inter-layer contracts via versioned Parquet files and strict execution lineage
- Engineered real-time acquisition for equity fundamentals & macroeconomic interest rates (FRED, Yahoo Finance) with automatic retry/backoff mechanism and run-level data provenance, giving every downstream dataset a traceable lineage
- Authored a no-look-ahead timing protocol using point-in-time timestamps (event/available/decision) to normalize daily panels powering a **Merton/KMV credit risk model**

Earnest Agriculture | *Marketing Analytics & Strategy Intern* | Champaign, IL

Jan 2026 – May 2026

- Boosted **organic traffic by 40%** and **sales by 14% YoY** by deploying a **Generative Engine Optimization (GEO)** strategy optimized for AI search engines (Perplexity, SearchGPT, Gemini)
- Developed a custom LLM-based evaluation pipeline (OpenAI API / Prompt Engineering) to automatically analyze, evaluate, and categorize competitor content patterns from a 1,500+ post dataset scraped via **Python/BeautifulSoup**
- Integrated **retrieval-augmented generation (RAG)** tools to streamline brand content creation aligned with **real-time Google Trends** and internal demand signals

SeeU International Inc. | *Data Analyst Intern* | Remote

Sep 2025 – Nov 2025

- Parsed 500K+ nested-JSON event records into **relational schemas** using **Python** and **SQL**, powering course-level KPI dashboards and funnel analytics
- Implemented **NLP data pre-processing** and automated text summarization workflows for user course feedback, improving data-driven curriculum iteration
- Designed an A/B test, drove a 15% lift in student engagement by translating funnel analysis into recommendations on course offerings and promotion timing

Walt Disney Company | *Operations & Analytics Intern* | Orlando, FL

Jun 2025 – Aug 2025

- Analyzed daily POS and attendance data covering 30,000+ visitors and \$150,000+ in daily sales to forecast demand across popcorn, bakery, and merchandise lines, optimizing replenishment schedules and reducing stockouts by 40%
- Built **inventory tracking dashboards** monitoring sell-through rates and shrinkage for data-driven reorders and promotional displays
- Trained 10 new hires on standardized POS procedures, **reducing ramp-up time by 25%** and improving transaction accuracy

AI/LLM RESEARCH

ScribeAR: Real-Time Diarization & Live Summarization | *Research Assistant*

Feb 2026 – Present

- Ship real-time speaker diarization for ScribeAR, an open-source live-captioning platform used by deaf and hard-of-hearing students, delivering stable speaker labels via overlap-voting reconciliation and word-level attribution
- Built a CPU-first streaming speaker diarization engine using **pyannote** (Sortformer-ready) and overlap-voting reconciliation for an open-source classroom accessibility platform
- Developed an LLM summarization pipeline that condenses lecture transcripts into real-time summaries for classroom accessibility
- Designed WCAG-AA compliant real-time UI components in React for dynamic speaker attribution and color rendering seamless accessibility standards.

Steer2Adapt for Efficient LLM Adaptation via Compositional Steering

Sep 2025 – Jan 2026

- Designed and implemented Steer2Adapt, a lightweight LLM adaptation framework that dynamically composed semantic steering vectors via Bayesian optimization, enabling efficient inference-time task adaptation without retraining
- Engineered a training-free LLM adaptation framework in **PyTorch** that dynamically composes semantic **steering vectors** via **Bayesian Optimization**, bypassing GPU-intensive fine-tuning and long-prompt context overhead
- Achieved an +8.2% avg accuracy lift across 9 tasks and 3 open-source LLMs (Llama-3.1, Qwen-2.5, Mistral) by building a reusable semantic concept subspace (e.g., reasoning and safety dimensions) and optimizing vector compositions from few-shot examples
- Developed a stability-aware objective function correcting wrong predictions, achieving robust and consistent gains across models
- Reduced adaptation cost by **eliminating fine-tuning and long prompts**, delivering low-latency control over LLM behavior