

Yuqing (Jacky) Jiang

Open to Relocate | (217)-898-8615 | jiangyuqing0508@outlook.com | [LinkedIn](#) | [GitHub](#)

EDUCATION

University of Illinois Urbana-Champaign

Bachelor of Science, Computer Science & Statistics (Dean's List)

Aug 2023 – May 2027

Cumulative GPA: 3.91/4.00

SKILLS & PUBLICATION

Machine Learning/AI: PyTorch, CUDA (GPU Programming), Scikit-learn, XGBoost, LLMs, PEFT, RAG, Prompt Engineering, LangChain, CrewAI

Programming & Analytics Tools: SQL, Power BI, Tableau, Python (Pandas, NumPy, Matplotlib), Cursor, Git, Java, C++

Publication: *Steer2Adapt: Dynamically Composing Steering Vectors for Efficient LLM Adaptation*. Pengrui Han, Xueqiang Xu, Keyang Xuan, Peiyang Song, Siru Ouyang, Runchu Tian, **Yuqing Jiang**, Cheng Qian, Pengcheng Jiang, Jiashuo Sun, Junxia Cui, Ming Zhong, Ge Liu, Jiawei Han, Jiaxuan You. **Accepted** at ICLR 2026 Re-Align Workshop; under review at NeurIPS 2026 (arXiv:2602.07276)

Personal Interests: Baseball, Tennis, Astronomy, Photography, Music

WORK EXPERIENCE

Pacific Financial Professional Association | *Quantitative Finance AI Intern* | Atlanta, GA (Remote)

Jun 2026 – Present

- Built a **four-layer market-data pipeline** (raw acquisition → cleaning → signal construction → publishing) for credit-risk analytics in Python/pandas, enforcing explicit inter-layer data contracts (versioned Parquet) and dependency rules between layers
- Engineered the acquisition layer ingesting equity fundamentals and interest rates (Yahoo Finance, FRED) with **retry/backoff and run-level provenance**, giving every downstream dataset a traceable lineage
- Normalized financial statements and corporate debt schedules into an analysis-ready daily panel powering a **Merton/KMV structural credit model**; authored a no-look-ahead timing protocol (event/available/decision timestamps) guaranteeing point-in-time correctness

Earnest Agriculture | *Marketing Analytics & Strategy Intern* | Champaign, IL

Jan 2026 – May 2026

- Boosted **organic blog traffic by 40%** (Google Analytics + Search Console) by pairing Webflow layout redesigns with a GEO content strategy targeting AI-search and LLM-referral queries
- Drove a **14% YoY increase in organic product sales** by mining Google Trends and internal sales data for demand signals and aligning the content calendar against them
- Scraped and structured a competitor-content dataset (Python + BeautifulSoup; 10+ brands, 1,500+ posts) and identified **top-performing content patterns** adopted into the posting strategy

SeeU International Inc. | *Data Analyst Intern* | Remote

Sep 2025 – Nov 2025

- Parsed nested-JSON event logs (**500K+ user-interaction records**) into analyzable tables with Python and SQL, powering course-level KPI reporting
- Built layered SQL views (window functions, CTEs) computing **10+ enrollment, engagement, and course-profitability KPIs** that directly informed curriculum and pricing decisions
- Drove a **15% lift in student engagement** by translating funnel analysis into recommendations on course offerings and promotion timing

Walt Disney Company | *Operations & Analytics Intern* | Orlando, FL

Jun 2025 – Aug 2025

- Analyzed daily POS and attendance data covering **30,000+ visitors and \$150,000+ in daily sales** to forecast demand across popcorn, bakery, and merchandise lines, optimizing replenishment schedules and **reducing stockouts by 40%**
- Built inventory tracking dashboards monitoring sell-through rates and shrinkage, enabling **data-driven reorders and promotional displays**
- Trained 10 new hires on standardized POS procedures, **reducing ramp-up time by 25%** and improving transaction accuracy

DP World PLC | *Data Analyst Intern* | Shanghai, China

May 2024 – Aug 2024

- Automated end-to-end data-entry and reporting pipelines with **Power Automate and UiPath**, cutting manual entry effort by 25% and improving data accuracy in shipment-tracking logistics
- Built a logistic regression model and **4 interactive Power BI dashboards** surfacing shipping-route and customer-region patterns for 50+ managers in contract-approval workflows; developed a Python web scraper to backfill missing client records
- Authored 2 user manuals and trained 20 staff across IT, Finance, and Operations, **driving adoption of the new RPA processes**

AI/LLM RESEARCH & PROJECT

ScribeAR: Real-Time Diarization & Live Summarization | *Undergraduate Research Assistant, UIUC*

Feb 2026 – Present

- Ship real-time **speaker diarization** for ScribeAR, an open-source live-captioning platform used by deaf and hard-of-hearing students, delivering stable speaker labels via overlap-voting reconciliation and word-level attribution
- Built a CPU-first streaming pipeline with pyannote (Sortformer-ready) and **WCAG-AA accessible speaker-color rendering** in the React client
- Developing an **LLM-based live-summarization pipeline** that condenses lecture transcripts into real-time summaries for classroom accessibility

Music-Streaming Analytics: Experimentation & Causal ML

Jun 2026 – Jul 2026

- Built an end-to-end user-behavior ML codebase on a simulated streaming dataset (**50K users, 1.5M events**) — churn/conversion models, a simulated A/B experiment, survival analysis, and uplift modeling across **10 CI-tested notebooks** and 75 unit tests
- Implemented **S-, T-, and X-learner uplift models** from scratch in scikit-learn and validated them against simulated ground-truth heterogeneous treatment effects; the guardrail-aware targeting policy captured **~70% of full-launch retention value with 38% of exposure**
- Designed a simulated **A/B test** end to end — power analysis, SRM gating, guardrail metrics, and peeking / multiple-testing failure demos — and made a **no-launch call** on a guardrail breach despite a significant +2.6pp retention lift
- Benchmarked tuned **XGBoost/LightGBM** against a logistic baseline with paired-bootstrap confidence intervals (no measurable lift) and fit a **Cox survival model** locating the week-1 churn-intervention window

AI Industry Monitor & LLM Benchmarking Platform

Mar 2026

- Developed a **full-stack platform** pairing a real-time global map and dashboard of AI-industry news, regulation, and market impact with an **LLM benchmarking arena** (Next.js, React, OpenAI API, RSS ingestion)
- Built a real-time data ingestion and alerting pipeline triggering **AI-generated reports** on market events and custom thresholds
- Designed an evaluation and benchmarking framework assessing **quantitative decision-making across frontier LLMs** (GPT, Gemini, Claude) by simulating trading strategies under dynamic market conditions
- Developed a real-time analytics interface with Recharts, exposing **equity curves, win rates, and per-decision model reasoning** for transparent cross-model comparison

Steer2Adapt for Efficient LLM Adaptation via Compositional Steering

Sep 2025 – Jan 2026

- Designed and implemented Steer2Adapt, a lightweight LLM adaptation framework that dynamically composed semantic steering vectors via Bayesian optimization, enabling efficient inference-time task adaptation without retraining
- Improved downstream performance by **+8.2% on average across 9 tasks and 3 open-source LLMs** (Llama-3.1, Qwen-2.5, Mistral) by building a reusable semantic concept subspace (e.g., reasoning and safety dimensions) and optimizing vector compositions from few-shot examples
- Developed a **stability-aware objective function** correcting wrong predictions, achieving robust and consistent gains across models
- Reduced adaptation cost by **eliminating fine-tuning and long prompts**, delivering low-latency control over LLM behavior
- Co-authored the paper — accepted at the **ICLR 2026 Workshop on Representational Alignment (Re-Align)**, under review at NeurIPS 2026 (arXiv:2602.07276)